

Matt Caraher

Dr. Forouraghi

Project 3

1 March 2022

Project 3

Part 1 MySQL: Importing Life_Expectancy.csv took about 10-15 minutes, and I started the SQL file by selecting my database and setting safe updates to 0. Running the first command, I deleted the rows that had a population equal to 0. There were 252 rows affected.

```
1 • use ids4db;  
2 • set SQL_SAFE_UPDATES = 0;  
3  
4 • delete from Life_Expectancy  
5   where population = 0;
```

To continue cleaning the data, I searched through each column and found that most of them had at least a few missing values and decided to replace them all with the average for each year. For example, I found the average life expectancy for the year 2013 and then replaced missing life expectancy values from the year 2013 with that average. I decided to find the averages dependent on year because I figured it would help to identify patterns in the data over the years. Here is the first column affected (life expectancy) and the year 2013, the same was done for each other column by each year respectively:

```

11 #Data cleaning
12 • select Country, Year
13 from Life_Expectancy
14 where Life_Expectancy = 0 and Year = 2013;
15
16 • select
17     @LifeExpectancy := avg(Life_Expectancy)
18 from Life_Expectancy
19 where Year = 2013;
20
21 • update Life_Expectancy
22 set Life_Expectancy = @Life_Expectancy
23 where Year = 2013 and Life_Expectancy = 0;
24
25

```

Output

#	Time	Action	Message
✓ 38	19:08:29	select Country, Year from Life_Expectancy where Life_...	2 row(s) returned
✓ 39	19:14:50	select Country, Year from Life_Expectancy where Life_...	2 row(s) returned
✓ 40	19:15:09	select @LifeExpectancy := avg(Life_Expectancy) from ...	1 row(s) returned
✓ 41	19:15:49	update Life_Expectancy set Life_Expectancy = @Life_...	2 row(s) affected Rows matched: 2 C

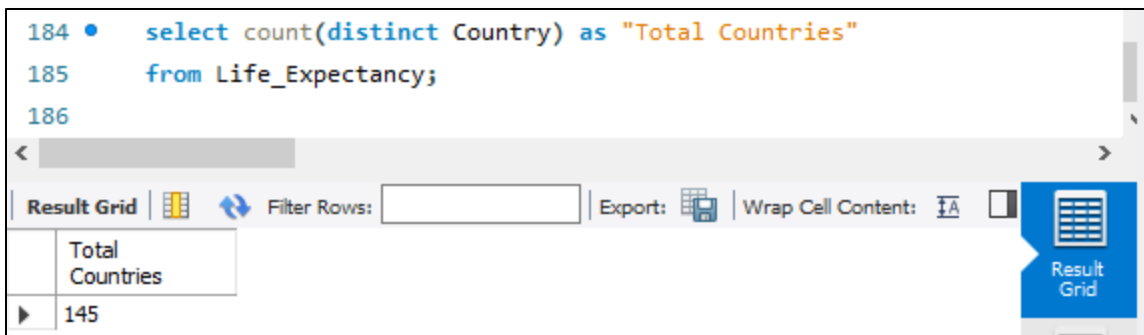
For Alcohol, Percentage Expenditure, and Total Expenditure, in year 2015, there were over 100 values equal to 0 in these rows, and so I had to specify that the avg variable is only taken from the values over 0. With the other columns and years, this was irrelevant because they each had 2 rows missing at most, so the effect of one 0 missing value was not problematic when I tested.

Here is an example:

```
154      #Alcohol 2015
155 •    select
156          @Alcohol2015 := avg(Alcohol)
157      from Life_Expectancy
158      where Year = 2015 and Alcohol > 0;
159
160 •    update Life_Expectancy
161      set Alcohol = @Alcohol2015
162      where Year = 2015 and Alcohol = 0;
163
```

To display the total count of countries after the cleaning, I used the following command to get 145 countries:

```
184 •    select count(distinct Country) as "Total Countries"
185      from Life_Expectancy;
186
```



Total Countries
145

To find the min and max of each mortality rate and list them by year in a table, I did the following (with the other year's commands that were cut off in the picture):

```
240 • select Year as "Year (first min, then max)", Country, Adult_Mortality
241 from Life_Expectancy
242 where Adult_Mortality = (select min(Adult_Mortality) from Life_Expectancy where Year = 2010) and Year = 2010
243 union all
244 select Year, Country, Adult_Mortality
245 from Life_Expectancy
246 where Adult_Mortality = (select max(Adult_Mortality) from Life_Expectancy where Year = 2010) and Year = 2010
247 union all
248 select Year, Country, Adult_Mortality
249 from Life_Expectancy
250 where Adult_Mortality = (select min(Adult_Mortality) from Life_Expectancy where Year = 2011) and Year = 2011
251 union all
252 select Year, Country, Adult_Mortality
253 from Life_Expectancy
```

Result Grid			
Filter Rows:		Export:	Wrap Cell Content: I
	Year (first min, then max)	Country	Adult_Mortality
►	2010	Fiji	2
	2010	Haiti	682
	2011	Italy	6
	2011	Zimbabwe	464
	2012	Senegal	2
	2012	Lesotho	513
	2013	Iceland	5
	2013	Lesotho	518
	2014	Kiribati	2
	2014	Lesotho	522
	2015	Tunisia	1
	2015	Lesotho	484

I did the same commands for the next tasks but replacing mortality rates with population, GDP, or schooling:

```
288
289 • select Year as "Year (first min, then max)", Country, Population
290 from Life_Expectancy
291 where Population = (select min(Population) from Life_Expectancy where Year = 2010) and Year = 2010
292 union all
293 select Year, Country, Population
294 from Life_Expectancy
295 where Population = (select max(Population) from Life_Expectancy where Year = 2010) and Year = 2010
296 union all
297 select Year, Country, Population
298 from Life_Expectancy
299 where Population = (select min(Population) from Life_Expectancy where Year = 2011) and Year = 2011
300 union all
301 select Year, Country, Population
302 from Life_Expectancy
303 where Population = (select max(Population) from Life_Expectancy where Year = 2011) and Year = 2011
304 union all
305 select Year, Country, Population
306 from Life_Expectancy
```

Result Grid | Filter Rows: | Export: | Wrap Cell Content:

	Year (first min, then max)	Country	Population
▶	2010	Hungary	123
	2010	Indonesia	242524123
	2011	Maldives	377
	2011	Brazil	198686688
	2012	Maldives	385
	2012	Indonesia	248883232
	2013	Palau	292
	2013	Pakistan	181712595
	2014	Maldives	41
	2014	India	1293859294
	2015	Sri Lanka	2966
	2015	Indonesia	258162113

```

338 • select Year as "Year (first min, then max)", Country, GDP
339 from Life_Expectancy
340 where GDP = (select min(GDP) from Life_Expectancy where Year = 2010) and Year = 2010
341 union all
342 select Year, Country, GDP
343 from Life_Expectancy
344 where GDP = (select max(GDP) from Life_Expectancy where Year = 2010) and Year = 2010
345 union all
346 select Year, Country, GDP
347 from Life_Expectancy
348 where GDP = (select min(GDP) from Life_Expectancy where Year = 2011) and Year = 2011
349 union all
350 select Year, Country, GDP
351 from Life_Expectancy
352 where GDP = (select max(GDP) from Life_Expectancy where Year = 2011) and Year = 2011
353 union all
354 select Year, Country, GDP
355 from Life_Expectancy

```

<



Result Grid   Filter Rows: Export:  Wrap Cell Content: 

	Year (first min, then max)	Country	GDP
▶	2010	Mauritius	8.376432
	2010	Norway	87646.75346
	2011	Senegal	18.25321
	2011	Luxembourg	115761.577
	2012	Guinea	52.3485646
	2012	Switzerland	83164.38795
	2013	Tajikistan	14.214412
	2013	Luxembourg	113751.85
	2014	Romania	12.27733
	2014	Luxembourg	119172.7418
	2015	Burundi	33.681223
	2015	Australia	56554.3876

```

387 • select Year as "Year (first min, then max)", Country, Schooling
388 from Life_Expectancy
389 where Schooling = (select min(Schooling) from Life_Expectancy where Year = 2010) and Year = 2010
390 union all
391 select Year, Country, Schooling
392 from Life_Expectancy
393 where Schooling = (select max(Schooling) from Life_Expectancy where Year = 2010) and Year = 2010
394 union all
395 select Year, Country, Schooling
396 from Life_Expectancy
397 where Schooling = (select min(Schooling) from Life_Expectancy where Year = 2011) and Year = 2011
398 union all
399 select Year, Country, Schooling
400 from Life_Expectancy
401 where Schooling = (select max(Schooling) from Life_Expectancy where Year = 2011) and Year = 2011
402 union all
403 select Year, Country, Schooling
404 from Life_Expectancy

```

Result Grid			
Filter Rows: <input type="text"/>			
Export:  Wrap Cell Content: 			
	Year (first min, then max)	Country	Schooling
▶	2010	Niger	4.5
	2010	Australia	19.5
	2011	Niger	4.8
	2011	Australia	19.8
	2012	South Sudan	4.9
	2012	Australia	20.1
	2013	South Sudan	4.9
	2013	Australia	20.3
	2014	South Sudan	4.9
	2014	Australia	20.4
	2015	South Sudan	4.9
	2015	Australia	20.4

For the Alcohol min-max table, there were a lot of repeated 0.01 values that tied for the min spot:

```
436 • select Year as "Year (first min, then max)", Country, Alcohol
437 from Life_Expectancy
438 where Alcohol = (select min(Alcohol) from Life_Expectancy where Year = 2010) and Year = 2010
439 union all
440 select Year, Country, Alcohol
441 from Life_Expectancy
442 where Alcohol = (select max(Alcohol) from Life_Expectancy where Year = 2010) and Year = 2010
443 union all
444 select Year, Country, Alcohol
445 from Life_Expectancy
446 where Alcohol = (select min(Alcohol) from Life_Expectancy where Year = 2011) and Year = 2011
447 union all
448 select Year, Country, Alcohol
449 from Life_Expectancy
450 where Alcohol = (select max(Alcohol) from Life_Expectancy where Year = 2011) and Year = 2011
451 union all
452 select Year, Country, Alcohol
453 from Life_Expectancy
454 where Alcohol = (select min(Alcohol) from Life_Expectancy where Year = 2012) and Year = 2012
```

<

Result Grid | Filter Rows: | Export: | Wrap Cell Content: [IA](#)

	Year (first min, then max)	Country	Alcohol
▶	2010	Afghanistan	0.01
	2010	Bangladesh	0.01
	2010	Mauritania	0.01
	2010	Estonia	14.97
	2011	Afghanistan	0.01
	2011	Bangladesh	0.01
	2011	Estonia	0.01
	2011	Fiji	0.01
	2011	Mauritania	0.01
	2011	Mongolia	0.01
	2011	Belarus	17.31
	2012	Afghanistan	0.01
	2012	Azerbaijan	0.01
	2012	Bangladesh	0.01

To decide if there is a correlation between population and life expectancy, I first tried sorting the countries by population and then by life expectancy. There was no clear trend, and so I edited the previous min-max chart for population and added the life expectancy column. With the following chart showing the min and max population from each year, the life expectancy was higher every year in the countries with smaller populations.

	Year (first min, then max) ▲	Country	Population	Life_Expectancy
	2010	Hungary	123	74.5
	2010	Indonesia	242524123	68.1
	2011	Maldives	377	77.3
	2011	Brazil	198686688	74.1
►	2012	Maldives	385	77.6
	2012	Indonesia	248883232	68.5
	2013	Palau	292	NULL
	2013	Pakistan	181712595	66
	2014	Maldives	41	78.2
	2014	India	1293859294	68
	2015	Sri Lanka	2966	74.9
	2015	Indonesia	258162113	69.1

Part 2 Python: The following picture is of my setup commands in Jupyter Notebooks. I had to install tabulate and the SQL connector, and the bottom of the picture includes the inputted column names.

```
[6]: import pandas as pd

[9]: pip install tabulate
Collecting tabulate
  Downloading tabulate-0.8.9-py3-none-any.whl (25 kB)
Installing collected packages: tabulate
Successfully installed tabulate-0.8.9
Note: you may need to restart the kernel to use updated packages.

[10]: from tabulate import tabulate

[11]: pip install mysql-connector-python
Requirement already satisfied: mysql-connector-python in c:\users\matt\anaconda3\lib\site-packages (8.0.28)
Requirement already satisfied: protobuf>=3.0.0 in c:\users\matt\anaconda3\lib\site-packages (from mysql-connector-python) (3.19.4)
Note: you may need to restart the kernel to use updated packages.

[12]: import mysql.connector as sql

[15]: db_connection = sql.connect(host='208.109.18.154', database='ids4db', user='ids4', password='Wst2394')
db_cursor = db_connection.cursor()
db_cursor.execute('SELECT * FROM Life_Expectancy')
table_rows = db_cursor.fetchall()
LifeExpDF = pd.DataFrame(table_rows)
LifeExpDF.columns = ['Country', 'Year', 'Life_Expectancy', 'Adult_Mortality', 'Alcohol', 'Percentage_Expenditure', 'BMI', 'Total_Expenditure', 'GDP', 'Population', 'Schooling']
print(tabulate(LifeExpDF, headers='keys', tablefmt='fancy_grid'))
```

	Country	Year	Life_Expectancy	Adult_Mortality	Alcohol	Percentage_Expenditure	BMI	Total_Expenditure	GDP	Population	Schooling
0	Afghanistan	2010	58.8	279	0.01	79.6794	16.7	9.2	553.329	2883167	9.2
1	Afghanistan	2011	59.2	275	0.01	7.09711	17.2	7.87	63.5372	2978599	9.5
2	Afghanistan	2012	59.5	272	0.01	78.1842	17.6	8.52	669.959	3696958	9.8
3	Afghanistan	2013	59.9	268	0.01	73.2192	18.1	8.13	631.745	31731688	9.9
4	Afghanistan	2014	59.9	271	0.01	73.5236	18.6	8.18	612.697	327582	10
5	Afghanistan	2015	65	263	0.01	71.2796	19.1	8.16	584.259	33736494	10.1
6	Albania	2010	76.2	91	5.28	41.8228	54.3	5.34	494.359	291321	12.5

I then looked at how adult mortality rates affect life expectancy. The Pearson r value being -0.69788 indicates that there is a somewhat strong negative correlation between adult mortality and life expectancy. As adult mortality goes up, life expectancy goes down.

```
# Calculate Pearson's correlation between adult mortality and life expectancy

from scipy.stats import pearsonr

data1 = LifeExpDF['Life_Expectancy']
data2 = LifeExpDF['Adult_Mortality']

corr, _ = pearsonr(data1, data2)
print('Pearsons correlation: %.5f' % corr)

Pearsons correlation: -0.69788
```

To test the correlation between life expectancy and eating habits, I compared life expectancy to BMI. With a Pearson's correlation of 0.4528, I know that these two have a weak but positive correlation.

```
In [21]: # Correlation between life expectancy and eating habits (BMI)

lifeExpectancyCol = LifeExpDF['Life_Expectancy']
bmiCol = LifeExpDF['BMI']

lifeExpectancyCol = np.nan_to_num(lifeExpectancyCol)

corr1, _ = pearsonr(lifeExpectancyCol, bmiCol)
print('Pearsons correlation : %.5f' % corr1)

Pearsons correlation : 0.45280
```

Life expectancy and alcohol have a weak positive correlation of 0.41422, like BMI. This number may be less or even negative if the data set included a greater number of extremes, in this case alcoholics. This is because alcoholics likely do not live as long as non-alcoholics.

```
In [24]: # Correlation between life expectancy and eating habits (BMI)

data1 = LifeExpDF['Life_Expectancy']
data2 = LifeExpDF['Alcohol']

data1 = np.nan_to_num(data1)

corr1, _ = pearsonr(data1, data2)
print('Pearsons correlation : %.5f' % corr1)

Pearsons correlation : 0.41422
```

To test the relation of a social factor and life expectancy, I chose schooling. This gave a strong positive correlation of 0.74597. This suggests that as the amount of schooling a sample has goes up, so does their life expectancy. Population had a correlation of -0.02695.

```
In [27]: # Correlation between life expectancy and schooling

data1 = LifeExpDF['Life_Expectancy']
data2 = LifeExpDF['Schooling']

data1 = np.nan_to_num(data1)

corr1, _ = pearsonr(data1, data2)
print('Pearsons correlation : %.5f' % corr1)

Pearsons correlation : 0.74597
```

The last relation I tested was economic factors, by comparing GDP to life expectancy. This gave a positive correlation of 0.44083. This is a very weak correlation much like alcohol an BMI to life expectancy.

```
In [29]: # Correlation between life expectancy and GDP

data1 = LifeExpDF['Life_Expectancy']
data2 = LifeExpDF['GDP']

data1 = np.nan_to_num(data1)

corr1, _ = pearsonr(data1, data2)
print('Pearsons correlation : %.5f' % corr1)

Pearsons correlation : 0.44083
```

I then compared total expenditure to life expectancy and got 0.17442, meaning that it has a very weak positive correlation. The strongest correlation seen is schooling to life expectancy, which previously showed that more schooling means a higher life expectancy.

The next step I took was to plot a linear regression model for life expectancy and schooling, to show the positive correlation. Below is also the line equation along with the R2 and MSE values:

```
In [77]: # Linear Regression between life expectancy and schooling
from sklearn.linear_model import LinearRegression

# Save Life expectancy in X
X = LifeExpDF.iloc[:, 2].values.reshape(-1,1)
X = np.nan_to_num(X)

# Save schooling in Y
Y = LifeExpDF.iloc[:, 10].values.reshape(-1,1)


# Create model
LRmodel = LinearRegression()

# Perform linear regression
LRmodel.fit(X, Y)

# Make predictions
Y_pred = LRmodel.predict(X)

# Visualize the dataset and the regression line:
plt.scatter(X, Y)
plt.plot(X, Y_pred, color='red')

plt.show()
```



```
In [90]: # Equation of the Line
print("The slope: ", LRmodel.coef_)
print("The intercept: ", LRmodel.intercept_)
print()

# Error analysis: MSE and R2
from sklearn.metrics import mean_squared_error, r2_score

print("MSE: ", mean_squared_error(Y, Y_pred))
print("R2: ", r2_score(Y, Y_pred))

The slope: [[0.23308557]]
The intercept: [-3.75920734]

MSE: 3.91800952000835
R2: 0.5564667277917688
```

Although BMI has a weak correlation, it is the next highest in the relations I tested. Here is a linear regression model along with R2 and MSE values:

```
In [78]: # Linear Regression between life expectancy and schooling
from sklearn.linear_model import LinearRegression

# Save life expectancy in X
X = LifeExpDF.iloc[:, 2].values.reshape(-1,1)
X = np.nan_to_num(X)

# Save schooling in Y
Y = LifeExpDF.iloc[:, 6].values.reshape(-1,1)


# Create model
LRmodel = LinearRegression()

# Perform linear regression
LRmodel.fit(X, Y)

# Make predictions
Y_pred = LRmodel.predict(X)

# Visualize the dataset and the regression line:
plt.scatter(X, Y)
plt.plot(X, Y_pred, color='red')

plt.show()
```



```
In [86]: # Equation of the Line
print("The slope: ", LRmodel.coef_)
print("The intercept: ", LRmodel.intercept_)
print()

# Error analysis: MSE and R2
from sklearn.metrics import mean_squared_error, r2_score

print("MSE: ", mean_squared_error(Y, Y_pred))
print("R2: ", r2_score(Y, Y_pred))

The slope: [[0.97937808]]
The intercept: [-29.05118008]

MSE: 336.4977643451747
R2: 0.20502985069323476
```

The last linear regression model I made was of life expectancy and alcohol, which has slightly less of a correlation than BMI. This chart, along with BMI, helped to target some of the outliers that still existed even after cleansing the data of 0's and null values. Below is the line equation along with the R2 and MSE values:

```
In [79]: # Linear Regression between Life expectancy and schooling
from sklearn.linear_model import LinearRegression

# Save Life expectancy in X
X = LifeExpDF.iloc[:, 2].values.reshape(-1,1)
X = np.nan_to_num(X)

# Save schooling in Y
Y = LifeExpDF.iloc[:, 4].values.reshape(-1,1)


# Create model
LRmodel = LinearRegression()

# Perform linear regression
LRmodel.fit(X, Y)

# Make predictions
Y_pred = LRmodel.predict(X)

# Visualize the dataset and the regression line:
plt.scatter(X, Y)
plt.plot(X, Y_pred, color='red')

plt.show()
```



```
In [84]: # Equation of the line
print("The slope: ", LRmodel.coef_)
print("The intercept: ", LRmodel.intercept_)
print()

# Error analysis: MSE and R2
from sklearn.metrics import mean_squared_error, r2_score

print("MSE: ", mean_squared_error(Y, Y_pred))
print("R2: ", r2_score(Y, Y_pred))

The slope: [[0.16523573]]
The intercept: [-7.45210736]

MSE: 11.927643414389264
R2: 0.1715750457910561
```

Out of the three models I created, I found that the first one, life expectancy and schooling, performs the best. This model had the highest R^2 value meaning it was the most accurate line relating to the data, and the MSE was the lowest meaning that the forecast line is the most accurate of the models.